

Terraform Resource Addressing

Resource Addressing

- A resource address is a unique way to target a specific resource
 - Acts as an identifier
 - Every resource in Terraform is assigned a unique resource address
 - What is a resource? Everything you create or define in Terraform, is a resource
- When to use a resource address?
 - In code - to create internal dependencies or get attributes from specific resources
 - Terraform CLI - running commands targeted to a specific resource
- A resource address consists of: `resource_type.resource_name[instance_index]`
 - `resource_type` - the type of resource, like an ec2 instance, an s3 bucket, depends completely on what Provider you are using (if using GCP provider, obviously the resource type is different)
 - `resource_name` - the name you've provided to that resource
 - `instance_index` - if the resource is created in an array or map (optional)
 - Not all resources follow this convention
 - Modules, variables, locals - are slightly different

Resource Addressing

- Example: We are using an AWS Provider, and creating an EC2 instance named "main"
 - Always consult documentation for resources available by Provider type
- `resource_type` = `aws_instance`
- `resource_name` = `main`
- `instance_index` = `None`, this is a standalone object and not part of any array or map
- To target this specific resource in your Workflow commands, the resource address is: `aws_instance.main`

```
resource "aws_instance" "main" {  
  provider = alias.west  
  ami = data.aws_ami.aws_linux_2.id  
  
  instance_type           = "t3.small"  
  associate_public_ip_address = true  
  
  vpc_security_group_ids = [aws_security_group.allow_ssh.id]  
  
  ebs_block_device {  
    device_name = "/dev/sdf"  
    volume_type = "gp2"  
    volume_size = "2"  
  }  
  
  tags = merge(local.tags, [  
    Name : "levelup_with_terraform"  
  ])  
}
```

Resource Addressing

```
data "aws_ami" "example" {  
  most_recent = true  
}
```

data.aws_ami.example

```
variable "image_id" {  
  type = string  
}
```

var.image_id

```
locals {  
  # Ids for multiple sets of  
  instance_ids = concat(aws  
}
```

local.instance_ids

```
module "vpc" {  
  source = "../modules/vpc"  
}
```

module.vpc

```
resource "aws_instance" "example" {  
  # ...  
}
```

aws_instance.example

Resource Addressing

```
module "vpc" {  
  source = "../modules/vpc"  
}
```

module.vpc.vpc_id

```
resource "aws_vpc" "main" {  
  cidr_block      = "10.0.0.0/16"  
  instance_tenancy = "default"  
  
  tags = {  
    Name = "main"  
  }  
}
```

```
output "vpc_id" {  
  value = aws_vpc.main.id  
}
```

module.vpc.aws_vpc.main

terraform plan -target module.vpc.aws_vpc.main

Resource Addressing

```
resource "azurerm_resource_group" "rg" {  
  for_each = {  
    a_group = "eastus"  
    another_group = "westus2"  
  }  
  name      = each.key  
  location = each.value  
}
```

azurerm_resource_group.rg["a_group"]

```
resource "aws_instance" "server" {  
  count = 4 # create four similar EC2 instances  
  
  ami          = "ami-a1b2c3d4"  
  instance_type = "t2.micro"  
}
```

aws_instance.server[0]

Why Is This Important?

- When writing Terraform code, resources may depend on other resources to be created first, or depend on their attributes
 - Example: Creating a subnet resource depends on your VPC resource id attribute

```
resource "aws_subnet" "main" {  
  vpc_id      = aws_vpc.main.id  
  cidr_block = "10.0.1.0/24"  
  
  tags = {  
    Name = "Main"  
  }  
}
```

Why Is This Important?

- When executing a terraform plan or terraform apply, you may want to run it for a particular resource, instead of the entire configuration
 - Example: `terraform apply -target aws_subnet.main`

```
resource "aws_subnet" "main" {  
  vpc_id      = aws_vpc.main.id  
  cidr_block = "10.0.1.0/24"  
  
  tags = {  
    Name = "Main"  
  }  
}
```