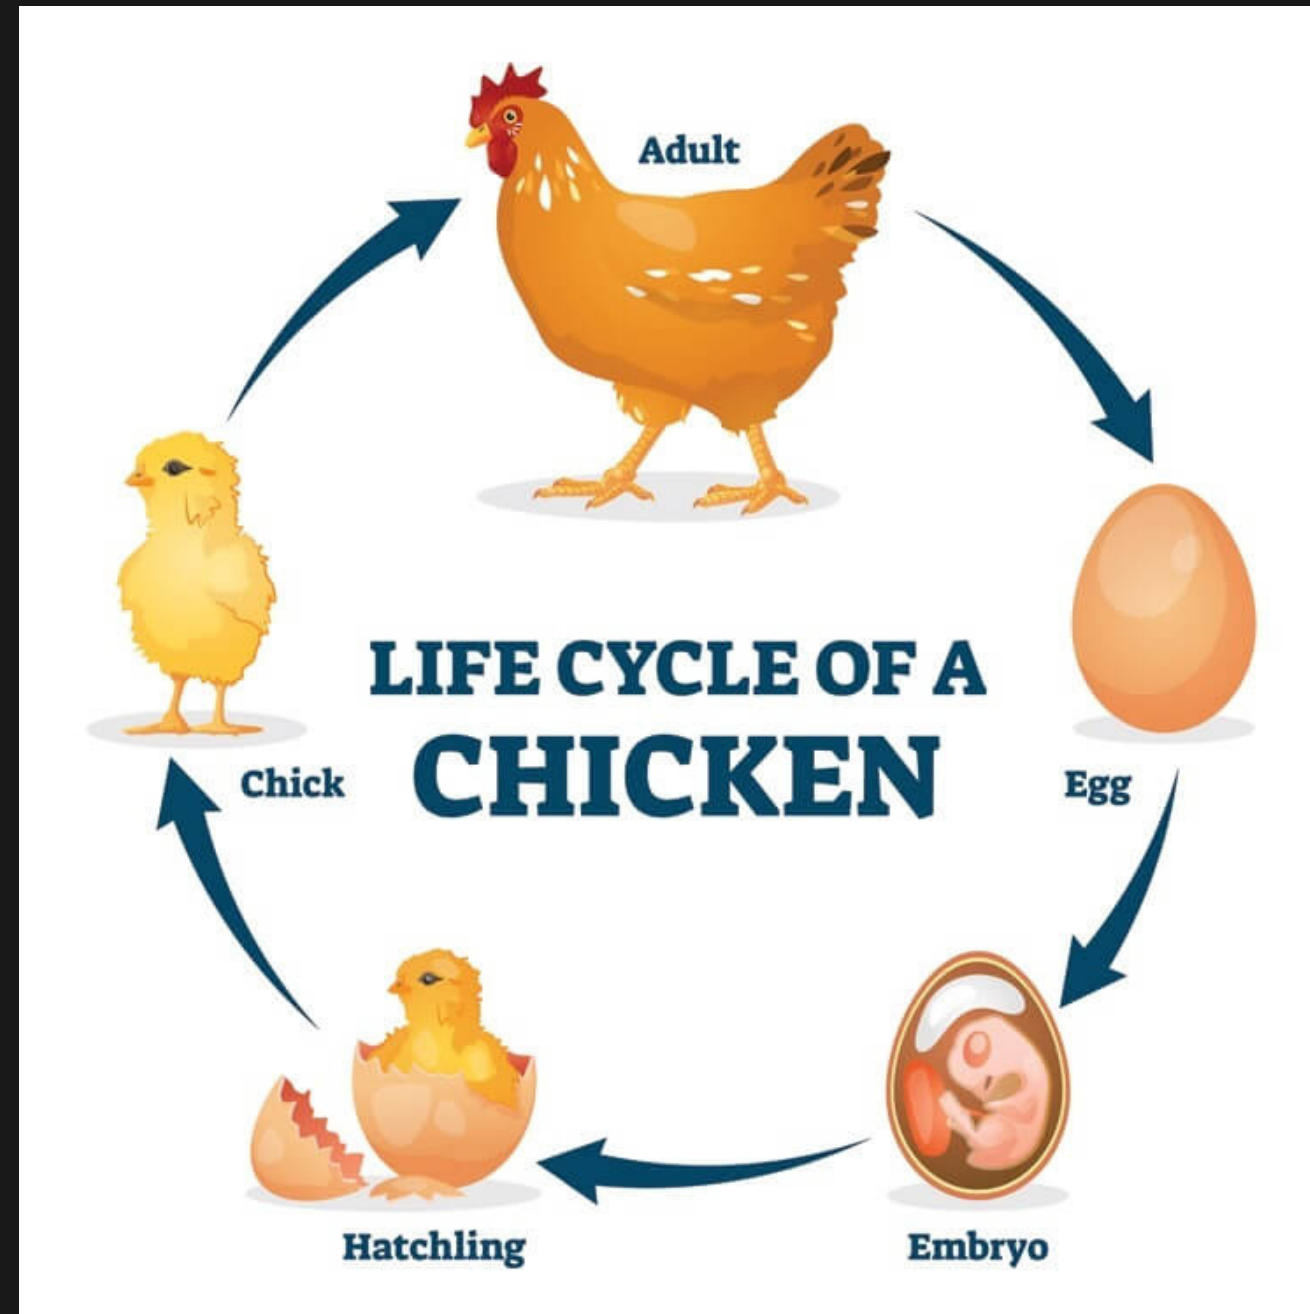


Terraform Resource Lifecycle



Resource Lifecycle

- A resource is defined in code with the `resource` block
 - Resource is created only if ran through `terraform apply`, and recorded record in state file
- Resource is recorded into state file
 - Allows Terraform to fully manage the resource
 - Creation, updating, destroying, grabbing metadata, etc.
 - Terraform consults state file when resource definition is altered, through code or manually
- Terraform makes proposed changes based on difference between code and state
- When resource is applied and created, it's attributes can be used by other Terraform configurations
 - Every resource exposes set of attributes
 - Create internal dependencies between resources by leveraging exposed attributes
 - Can use `data source` block to get values from AWS into Terraform

Lifecycle Arguments

- Terraform provides global `lifecycle` block which alters the behavior of Terraform to that resource
- The lifecycle nested block is available for every resource
 - Not Provider specific, every resource has standard lifecycle behavior
 - This is a nested block, meaning it must be defined inside a resource

```
resource "azurerm_resource_group" "example" {  
  # ...  
  
  lifecycle {  
    create_before_destroy = true  
  }  
}
```

Lifecycle Arguments

- There are four available arguments for the lifecycle block
 - `create_before_destroy`
 - `prevent_destroy`
 - `ignore_changes`
 - `replace_triggered_by`
- The most commonly used arguments are `create_before_destroy`, `prevent_destroy`, `ignore_changes`
- Lifecycle arguments should be used sparingly!
 - Alters the default behavior of how resource is created and updated
 - Can lead to many unintended side effects, potentially corrupting state file
 - `ignore_changes` is safe
 - `create_before_destroy`, `prevent_destroy` has significant impact

create_before_destroy

- Accepts a boolean value (true or false)
- By default, if Terraform detects a change and must update a resource, it will first destroy the resource if the update cannot be made in-place
- If `create_before_destroy` is enabled, then Terraform reverses that behavior
 - A new resource gets created, with your proposed changes, then the old resource is destroyed
- Useful if the resource must be highly available, minimize downtime, or must maintain some kind of state
- Sounds like a no-brainer to enable, but you must fully understand potential side effects
 - Difficult to fully understand all side effects
 - Can easily corrupt state file depending on resource type

prevent_destroy

- Accepts a boolean value (true or false)
- By default, if Terraform detects a change and must update a resource, it will first destroy the resource if the update cannot be made in-place
- If `prevent_destroy` is enabled, then Terraform will reject any plans that proposes the destruction of the resource, and provide an error
- This setting protects the resource from accidentally being deleted
- Useful for sensitive resources which under no circumstances, must be removed
 - State buckets, Cognito user pools, Highly available load balancers, etc.
- Similar to `create_before_destroy`, think about ramifications first

ignore_changes

- Accepts a list of attribute names
- When Terraform detects change in a resource argument, it will propose an in-place update, or a full re-creation of the resource
- If argument/attribute is supplied to `ignore_changes`, Terraform will ignore request to update the resource
- Commonly used when refactoring / changing existing modules
 - You don't want existing infrastructure to be affected by the changes
- Some resources get updated by CI/CD, Lambda is common example

```
resource "aws_instance" "example" {  
  # ...  
  
  lifecycle {  
    ignore_changes = [  
      # Ignore changes to tags, e.g. because a management agent  
      # updates these based on some ruleset managed elsewhere.  
      tags,  
    ]  
  }  
}
```


ignore_changes

- If resource has yet to be created or is destroyed, Terraform will create resource like normal
 - The `ignore_changes` lifecycle is ignored during the creation phase
- Only the attributes defined by the resource block are valid
 - As always, visit the Terraform documentation to get list of all available attributes

```
resource "aws_instance" "example" {  
  # ...  
  
  lifecycle {  
    ignore_changes = [  
      # Ignore changes to tags, e.g. because a management agent  
      # updates these based on some ruleset managed elsewhere.  
      tags,  
    ]  
  }  
}
```